

C++ Best Practices

45ish Simple Rules with Specific
Action Items for Better C++

Full **COLOR** Syntax Highlighting

Jason Turner



C++ Best Practices

45ish Simple Rules with Specific Action Items for Better C++

Jason Turner

This book is for sale at <http://leanpub.com/cppbestpractices>

This version was published on 2021-01-14



This is a [Leanpub](#) book. Leanpub empowers authors and publishers with the Lean Publishing process. [Lean Publishing](#) is the act of publishing an in-progress ebook using lightweight tools and many iterations to get reader feedback, pivot until you have the right book and build traction once you do.

© 2020 - 2021 Jason Turner

For my wife Jen, and her love of silly alpaca.

*The cover picture was taken by my wife while visiting Red Fox Alpaca Ranch in
Evergreen, Colorado.*

Contents

1. Introduction	1
2. About Best Practices	2
3. Use the Tools: Automated Tests	5
4. Use the Tools: Continuous Builds	7
5. Use the Tools: Compiler Warnings	9
6. Exercise: Use the Tools: Static Analysis	11
7. Use the Tools: Sanitizers	12
8. Slow Down	14
9. C++ Is Not Magic	15
10. C++ Is Not Object-Oriented	17
11. Learn Another Language	19
12. <code>const</code> Everything That's Not <code>constexpr</code>	21
13. <code>constexpr</code> Everything Known at Compile Time	23
14. Prefer <code>auto</code> in Many Cases.	26
15. Prefer ranged-for Loop Syntax Over Old Loops	32
16. Use <code>auto</code> in ranged for loops	34

CONTENTS

17. Prefer Algorithms Over Loops	36
18. Don't Be Afraid of Templates	38
19. Don't Copy and Paste Code	40
20. Follow the Rule of 0	42
21. If You Must Do Manual Resource Management, Follow the Rule of 5	44
22. Don't Invoke Undefined Behavior	47
23. Never Test for <code>this To Be nullptr</code>, It's UB	49
24. Never Test for A Reference To Be <code>nullptr</code>, It's UB	52
25. Avoid <code>default</code> In <code>switch</code> Statements	54
26. Prefer Scoped <code>enums</code>	58
27. Prefer <code>if constexpr</code> over <code>SFINAE</code>	61
28. Constrain Your Template Parameters With Concepts (C++20)	64
29. De-template-ize Your Generic Code	68
30. Use Lippincott Functions	71
31. Be Afraid of Global State	74
32. Make your interfaces hard to use wrong.	75
33. Consider If Using the API Wrong Invokes Undefined Behavior	76
34. Use <code>[[nodiscard]]</code> Liberally	78
35. Use Stronger Types	81
36. Don't return raw pointers	85
37. Prefer Stack Over Heap	86

CONTENTS

38. No More new!	89
39. Know Your Containers	91
40. Avoid <code>std::bind</code> and <code>std::function</code>	93
41. Skip C++11	97
42. Don't Use <code>initializer_list</code> For Non-Trivial Types	100
43. Use the Tools: Build Generators	102
44. Use the Tools: Package Managers	105
45. Improving Build Time	106
46. Use the Tools: Multiple Compilers	108
47. Fuzzing and Mutating	110
48. Continue Your C++ Education	115
49. Thank You	118
50. Bonus: Understand The Lambda	120

1. Introduction

My goal as a trainer and a contractor (seems to be) to work me out of a job. I want everyone to:

1. Learn how to experiment for themselves
2. Not just believe me, but test it
3. Learn how the language works
4. Stop making the same mistakes of the last generation

I'm thinking about changing my title from "C++ Trainer" to "C++ Guide." I always adapt my courses and material to the class I currently have. We might agree on X, but I change it to Y halfway through the first day to meet the organization's needs.

Along the way, we experiment and learn as a group. I often also learn while teaching. Every group is unique; every class has new questions.

But a lot of the questions are still the same ones over and over (to the point where I get to look like a mind reader, that bit's fun

Hence, this book (and the twitter thread that it came from) to spread the word on the long-standing best practices.

I wrote the book I wanted to read. It's intentionally straightforward, short, to the point, and has specific action items.

2. About Best Practices

Best Practices, quite simply, are about

1. Reducing common mistakes
2. Finding errors quickly
3. Without sacrificing (and often improving) performance

Why Best Practices?

First and foremost, let's get this out of the way:

Your Project Is Not Special

If you are programming in C++, you, or someone at your company, cares about performance. Otherwise, they'd probably be using some other programming language. I've been to many companies who all tell me they are special because they need to do things fast!

Spoiler alert: they are all making the same decisions for the same reasons.

There are very few exceptions. The outliers who make different decisions: they are the organizations that are already following the advice in this book.

What's The Worst Than Can Happen?

I don't want to be depressing, but let's take a moment to ponder the worst-case scenario if your project has a critical flaw.

Game

Serious flaws lead to remote vulnerability or attack vector.

Financial

Serious flaws lead to [large amounts of lost money, accelerating trades, market crash](#)¹.

Aerospace

Serious flaws lead to lost spacecraft or [human life](#)².

Your Industry

Serious flaws lead to... Lost money? Lost jobs? Remote hacks? Worse?

Examples

Examples throughout this book use `struct` instead of `class`. The only difference between `struct` and `class` is that `struct` has all members by default `public`. Using `struct` makes examples shorter and easier to read.

Exercises

Each section has one or more exercises. Most do not have a right or wrong answer.



Exercise: Look for exercises

Throughout the following chapters, you'll see exercises like this one. Look for them!

Exercises are:

- Practical, and apply to your current code base to see immediate value.

¹https://en.wikipedia.org/wiki/2010_flash_crash

²<https://spectrum.ieee.org/aerospace/aviation/how-the-boeing-737-max-disaster-looks-to-a-software-developer>

- Make you think and understand the language a little bit deeper by doing your own research.

Links and References

I've made an effort to reference those who I learned from and link to their talks where possible. If I've missed something, please let me know.

3. Use the Tools: Automated Tests

You need a single command to run tests.
If you don't have that, no one will run the tests.

- [Catch2](#)¹ - popular and well supported testing framework from [Phil Nash](#)² and [Martin Hořeňovský](#)³
- [doctest](#)⁴ - similar to catch2, but trimmed for compile-time performance
- [Google Test](#)⁵
- [Boost.Test](#)⁶ - testing framework, boost style.

[ctest](#)⁷ is a test runner for CMake that can be used with any of the above frameworks. It is utilized via the [add_test](#)⁸ feature of CMake.

You need to be familiar with these tools, what they do, and pick from them.

Without automated tests, the rest of this book is pointless. You cannot apply the practical exercises if you cannot verify that you did not break the existing code.

Oleg Rabaev on CppCast stated:

- If a component is hard to test, it is not properly designed.
- If a component is easy to test, it is a good indication that it is properly designed.

¹<https://github.com/catchorg/Catch2>

²https://twitter.com/phil_nash

³https://twitter.com/horenmar_ctu

⁴<https://github.com/onqtam/doctest>

⁵<https://github.com/google/googletest>

⁶https://www.boost.org/doc/libs/1_74_0/libs/test/doc/html/index.html

⁷<https://cmake.org/cmake/help/latest/manual/ctest.1.html>

⁸https://cmake.org/cmake/help/latest/command/add_test.html

- If a component is properly designed, it is easy to test.



Exercise: Can you run a single command to run a suite of tests?

- Yes: Excellent! Run the tests and make sure they all pass!
- No: Does your program produce output?
 - Yes: Start with “[Approval Tests](#)”⁹, which will give you the foundation you need to start started with testing.
 - No: Develop a strategy for how to implement some minimal form of testing.

Resources

- CppCon 2018: Phil Nash “Modern C++ Testing with Catch2”¹⁰
- CppCon 2019: Clare Macrae “Quickly Testing Legacy C++ Code with Approval Tests”¹¹
- C++ on Sea 2020: Clare Macrae “Quickly and Effectively Testing Legacy C++ Code with Approval Tests”¹²

⁹<https://cppcast.com/clare-macrae/>

¹⁰https://youtu.be/Ob5_XZrFQH0

¹¹<https://youtu.be/3GZHvcdq32s>

¹²https://youtu.be/tXEuf_3VzRE

4. Use the Tools: Continuous Builds

Without automated tests, it is impossible to maintain project quality.

In the C++ projects I have worked on throughout my career, I've had to support some combination of:

- x86
- x64
- SPARC
- ARM
- MIPSSEL

On

- Windows
- Solaris
- MacOS
- Linux

When you start to combine multiple compilers across multiple platforms and architectures, it becomes increasingly likely that a significant change on one platform will break one or more other platforms.

To solve this problem, enable continuous builds with continuous tests for your projects.

- Test all possible combinations of platforms that you support

- Test Debug and Release separately
- Test all configuration options
- Test against newer compilers than you support or require



If you don't require 100% tests passing, you will never know the code's state.



Exercise: Enable continuous builds

Understand your organization's current continuous build environment. If one does not exist, what are the barriers to getting it set up? How hard would it be to get something like GitLab, GitHub actions, Appveyor, or Travis set up for your projects?

5. Use the Tools: Compiler Warnings

There are many warnings you are not using, most of them beneficial. `-Wall` is *not* all warnings on GCC and Clang. `-Wextra` is still barely scratching the surface!



`/Wall` on MSVC is *all* of the warnings. Our compiler writers do not recommend using `/Wall` on MSVC or `-Weverything` on Clang, because many of these are diagnostic warnings. GCC does not provide an equivalent.

Strongly consider `-Wpedantic` (GCC/Clang) and `/permissive-` (MSVC). These command line options disable language extensions and get you closer to the C++ standard. The more warnings you enable today, the easier time you will have with porting to another platform in the future.



Exercise: Enable More Warnings

1. Explore the set of warnings available with your compiler. Enable as many as you can.
2. Fix the new warnings generated.
3. Goto 1.



MSVC has an excellent set of warnings that can be enabled by warning level. You can start with `/W1` and work your way up to `/W4` as you fix each set of warnings.

This process will feel tedious and meaningless, but these warnings will catch real bugs.

Resources

- [C++ Best Practices website curated list of warnings](#)¹
- [GCC's full warning list](#)²
- [Clang's full warning list](#)³
- [MSVC's Compiler warnings that are off by default](#)⁴
- [C++ Weekly Ep 168 - Discovering Warnings You Should Be Using](#)⁵

¹https://github.com/lefticus/cppbestpractices/blob/master/02-Use_the_Tools_Available.md#compilers

²<https://gcc.gnu.org/onlinedocs/gcc/Warning-Options.html>

³<https://clang.llvm.org/docs/DiagnosticsReference.html>

⁴<https://docs.microsoft.com/en-us/cpp/preprocessor/compiler-warnings-that-are-off-by-default?view=vs-2019>

⁵<https://youtu.be/IOo8gTDMFkM>

6. Exercise: Use the Tools: Static Analysis

Static analysis tools are tools that analyze your code without compiling or executing it. Your compiler is one such tool and your first line of defense.

Many such tools are free and some are free for open source projects.

cppcheck and clang-tidy are two popular and free tools with major IDE and editor integration.



Enable More Static Analysis

Visual Studio: look into Microsoft's static analyzer that ships with it. Consider using Clang Power Tools. Download cppcheck's addon for visual studio

CMake: Enable cppcheck and clang-tidy integration

Resources

- [cppbestpractices.com](https://github.com/lefticus/cppbestpractices/blob/master/02-Use_the_Tools_Available.md#static-analyzers) list of static analyzers¹

¹https://github.com/lefticus/cppbestpractices/blob/master/02-Use_the_Tools_Available.md#static-analyzers

7. Use the Tools: Sanitizers

The sanitizers are runtime analysis tools for C++ and are built into GCC, Clang, and MSVC.

If you are familiar with Valgrind, the sanitizers provide similar functionality but many orders of magnitude faster than Valgrind.

Available sanitizers are:

- Address (ASan)
- Undefined Behavior (UBSan) (More on Undefined Behavior later)
- Thread
- DataFlow (use for code analysis, not finding bugs)
- Lib Fuzzer (addressed in a later chapter)

Address sanitizer, UB Sanitizer, Thread sanitizer can find many issues almost like magic. Support is currently increasing in MSVC at the time of this book's writing, while GCC and Clang have more established support for the sanitizers.

[John Regehr](https://twitter.com/johnregehr)¹ recommends always enabling ASan and UBSan during development.

When an error such as an out of bounds memory access occurs, the sanitizer will give you a report of what conditions led to the failure, often with suggestions for fixing the problem.

¹<https://twitter.com/johnregehr>

You can enable Address and Undefined Behavior sanitizers with a command similar to:

```
gcc -fsanitize=address,undefined <filetocompile>
```

Sanitizers must also be enabled during the linking phase of the project build.



Examples for how to use sanitizers with CMake exist in the [C++ Starter Project](#)²



Exercise: Enable Sanitizers

- Investigate how to add sanitizer support for your existing project
- Enable ASan first
- Run the full test suite and investigate any problems found
- Enable UBSan second
- Run full test suite again

End goal: get all tests running with ASan, and UBSan enabled on your continuous build environment.

Resources

- [AddressSanitizer \(ASan\) for Windows with MSVC](#)³
- [Sanitizers source and documentation on GitHub](#)⁴
- [Clang AddressSanitizer documentation](#)⁵
- [Clang UndefinedBehaviorSanitizer documentation](#)⁶

²https://github.com/lefticus/cpp_starter_project

³<https://devblogs.microsoft.com/cppblog/addresssanitizer-asan-for-windows-with-msvc/>

⁴<https://github.com/google/sanitizers>

⁵<https://clang.llvm.org/docs/AddressSanitizer.html>

⁶<https://clang.llvm.org/docs/UndefinedBehaviorSanitizer.html>

8. Slow Down

Dozens of solutions exist in C++ for any given problem. Dozens of more opinions exist for which of these solutions are the best. Copying and pasting from one application to another is easy. Forging ahead with the solutions you are comfortable with is easy.

How many times have you said, “wow, this is going to take a complicated class hierarchy to implement this solution?” Or what about “I guess I need to add macros here to implement these common functions.”

- If the solution seems large or complex, stop.
- Now is a good time to go for a walk and ponder the solution.
- When you’re done with your walk, discuss the design with a coworker, pet, or [rubber duck](https://rubberduckdebugging.com/)¹.

Still haven’t found a more straightforward solution you are happy with? Ask on Twitter or Slack if you can.

The key point is to not forge ahead blindly with the solutions with which you are comfortable. Be willing to stop for a minute. The older I get, the less time I spend programming, and the more time I spend thinking. In the end, I implement the solution as fast or faster than I used to and with less complexity.

¹<https://rubberduckdebugging.com/>

9. C++ Is Not Magic

This section is just a reminder that we can reason about all aspects of C++. It's not a black box, and it's not magic.

If you have a question, it's usually easy to construct an experiment that helps you answer the question for yourself.

A favorite tool of mine is this simple class that prints a debug message whenever a special member function is called.

Understanding object lifetime tool

```
#include <cstdio>

struct S {
    S(){ puts("S()"); }
    S(const S &){ puts("S(const S &)"); }
    S(S &&){ puts("S(S &&)"); }
    S &operator=(const S &){
        puts("operator=(const S &)");
        return *this;
    }
    S &operator=(S &&){
        puts("operator=(S &&)");
        return *this;
    }
    ~S() { puts("~S()"); }
};
```



Exercise: Build your first C++ experiment.

Do you have a question about C++ that's been nagging you? Can you design an experiment to test it? Remember that Compiler Explorer now allows you to execute code.



Exercise: Start collecting your experiments.

Once you have created an experiment and test, be sure to save it. Consider using GitHub gists as a simple way to save and share your tests with others.

Resources

- [A quick start example with Compiler Explorer.](https://godbolt.org/z/3eGP56)¹

¹<https://godbolt.org/z/3eGP56>

10. C++ Is Not Object-Oriented

Bjarne Stroustrup in The C++ Programming Language 3rd Edition states:

C++ is a general-purpose programming language with a bias towards systems programming that

- is a better C,
- supports data abstraction,
- supports object-oriented programming, and
- supports generic programming.

You must understand that C++ is a multi-discipline programming language to make the most of the language. C++ supports effectively all of the programming paradigms that exist today.

- Procedural
- Functional
- Object-Oriented
- Generic
- Compile-Time (constexpr and template metaprogramming)

Knowing when it is appropriate to use each of these tools is the key to writing good C++. Projects that rigidly stick to one paradigm miss out on the best features of the language.



Don't try to use every technique possible all of the time. You will end up with a mess of difficult to maintain and read code. Appropriately using the appropriate techniques at the appropriate times takes discipline and practice.



Exercise: Question your current design.

If you could break out of the current design your project is using, what would you do differently?

Resources

- [Functional Programming in C++¹](https://www.manning.com/books/functional-programming-in-c-plus-plus?a_aid=FPinCXX&a_bid=441f12cc)
- [C++ Weekly Ep 137: C++ Is Not an Object Oriented Language²](https://youtu.be/AUT201AXeJg)

¹https://www.manning.com/books/functional-programming-in-c-plus-plus?a_aid=FPinCXX&a_bid=441f12cc

²<https://youtu.be/AUT201AXeJg>

11. Learn Another Language

Considering that [C++ is not an object-oriented language](#), you have to know many different techniques to make the most of C++.

The following exercises will help expose you to other languages. But the fact is that currently, few languages are pure single paradigm languages.

Every language has its preferred way of doing things that work within the language's preferred paradigm.

Ben Deane [recommends this set of languages that all programmers should learn](#)¹:

- ALGOL family (C and descendants)
- Forth
- Lisp and dialects
- Haskell
- Smalltalk
- Erlang



Exercise: Pick a functional language to learn

Can you find a pure functional language?

¹<http://www.elbeno.com/blog/?p=420>



Exercise: Pick an object-oriented language to learn

Finding a pure object-oriented language is even harder! Even Java has lambda functions these days.



Exercise: Pick a language with a different syntax

Languages that look like C-family languages will likely be more comfortable for you. Try to find a language that looks different and stretches your mind.

Resources

- [Execution in the Kindom of Nouns²](https://steve-yegge.blogspot.com/2006/03/execution-in-kingdom-of-nouns.html) - gets you thinking about different programming paradigms

²<https://steve-yegge.blogspot.com/2006/03/execution-in-kingdom-of-nouns.html>

12. `const` Everything That's Not `constexpr`

Many people (like Kate Gregory and James McNellis) have said this many times. Making objects `const` does two things:

1. It forces us to think about the initialization and lifetime of objects, which affects performance.
2. Communicates meaning to the readers of our code.

And as an aside, if it's a static object, the compiler is now free to move it into the constants portion of the binary, which can affect the optimizer.



Exercise: Look for `const` opportunities.

As you read through your code, you should look for variables that are not `const` and make them `const`.

- If a variable is not `const`, ask why not?
- Would using a lambda or adding a named function allow you to make the value `const`?

Using a lambda to initialize a const object.

```
const auto data = [](){ // no parameters
    std::vector<int> result;
    // fill result with things.
    return result;
}(); // immediately invoked
```



Because of RVO, using a lambda will likely not add any overhead and may increase performance.

Did you make any static variables const in the process?
Then [go to the constexpr exercise](#).



You probably don't want to make class members const; it can break essential things, and sometimes silently.

Resources

- [CppCon 2014: James McNellis & Kate Gregory “Modernizing Legacy C++ Code”](#)¹
- [CppCon 2019: Jason Turner “C++ Code Smells”](#)²
- [The implication of const or reference member variables in C++](#)³
- [C++Now 2018: Ben Deane “Easy to Use, Hard to Misuse: Declarative Style in C++”](#)⁴ (Builds on techniques that make applying const easier.)

¹<https://youtu.be/LDxAgMe6D18>

²https://youtu.be/f_tLQl0wLUM

³<https://lesleylai.info/en/const-and-reference-member-variables/>

⁴<https://youtu.be/2ouxETt75R4>

13. constexpr Everything Known at Compile Time

Gone are the days of `#define`. `constexpr` should be your new default! Unfortunately, people over-complicate `constexpr`, so let's break down the simplest thing.

If you see something like (I've seen in real code):

static const data known at compile time.

```
static const std::vector<int> angles{-90,-45,0,45,90};
```

This really needs to be:

Moving static const to static constexpr.

```
static constexpr std::array<int, 5> angles{-90,-45,0,45,90};
```



`static constexpr` here is necessary to make sure the object is not reinitialized each time the function / declaration is encountered. With `static` the variable lasts for the lifetime of the program, and we know it will be initialized exactly once.

The difference is threefold.

- The size of the array is now known at compile time
- We've removed dynamic allocations
- We no longer pay the cost of accessing a static

The main gains come from the first two, but we need a constexpr mindset to be looking for this kind of opportunity. We also need constexpr knowledge to see how to apply it in the more complex cases.

The difference can be significant.



Exercise: constexpr Your const Values

While reading code, look at all `const` values. Ask, “is this value known at compile time?” If it is, what would it take to make the value `constexpr`?



Exercise: static constexpr Your static const Values

Go through your current code base and look for code that is currently `static const`. You probably have something, somewhere.

- If it’s currently `static const`, it’s likely the size and data are known at compile time.
- Can this code become `constexpr`?
- What is preventing it from being `constexpr`?
- How much work would it take to modify the functions populating the `static const` data so that they are also `constexpr`?

Resources

- [C++Now 2017: Ben Deane & Jason Turner “constexpr ALL the things”¹](#) (a bit out of date with modern `constexpr` techniques)
- [C++ Weekly Ep 233: constexpr map vs std::map²](#)

¹<https://youtu.be/HMB9oXFobJc>

²<https://youtu.be/INn3xa4pMfg>

- Meeting C++ 2017: Jason Turner “Practical constexpr”³
- C++ Russia 2019: Hana Dusíková “A state of compile time regular expressions”⁴

³<https://youtu.be/xtf9qkDTrZE>

⁴https://youtu.be/r_ZASJFQGQI

14. Prefer auto in Many Cases.

I'm not an [Almost Always Auto](#) (AAA) person, but let me ask you this: What is the result type of `std::count`?

My answer is, "I don't care."

```
const auto
```

```
const auto result = std::count( /* stuff */ );
```

or, if you prefer:

```
auto const
```

```
auto const result = std::count( /* stuff */ );
```

Using `auto` avoids unnecessary conversions and data loss. Same as ranged-for loops. `auto` requires initialization, the same as `const`, the same reasoning for why that's good.

Example:

Possible expensive conversion.

```
const std::string value = get_string_value();
```

What is the return type of `get_string_value()`? If it is `std::string_view` or `const char *`, we will get a potentially costly conversion on all compilers with no diagnostic.

No possible expensive conversion.

```
// avoids conversion
const auto value = get_string_value();
```

Furthermore, auto return types actually can significantly simplify generic code.

C++ 98 template usage.

```
// our example from "Don't Be Afraid of Templates"
template<typename Arithmetic>
Arithmetic divide(Arithmetic numerator, Arithmetic denominator) {
    return numerator / denominator;
}
```

This code forces us to use the same type for both the numerator and denominator (play with this and see the weird compile errors you get).

C++ 98 template made more generic?

```
template<typename Numerator, typename Denominator>
/*what's the return type*/
divide(Numerator numerator, Denominator denominator) {
    return numerator / denominator;
}
```

C++98 provides no solution to this problem, but C++11 does.

C++11 trailing return types.

```
// use trailing return type
template<typename Numerator, typename Denominator>
auto divide(Numerator numerator, Denominator denominator)
    -> decltype(numerator / denominator)
{
    return numerator / denominator;
}
```

But in C++14, we can leave off the return type altogether (remember to [Skip C++11](#)).

C++14 auto return types.

```
template<typename Numerator, typename Denominator>
auto divide(Numerator numerator, Denominator denominator)
{
    return numerator / denominator;
}
```



Exercise: Become familiar with auto deduction.

Ex1: what is the type of val?

```
const int *get();
```

```
int main() {  
    const auto val = get();  
}
```

Ex2: what is the type of val?

```
const int &get();
```

```
int main() {  
    const auto val = get();  
}
```

Ex3: what is the type of val?

```
const int *get();
```

```
int main() {  
    const auto *val = get();  
}
```

Ex4: what is the type of val?

```
const int &get();

int main() {
    const auto &val = get();
}
```

Ex5: what is the type of val?

```
const int *get();

int main() {
    const auto &val = get();
}
```

Ex6: what is the type of val?

```
const int &get();

int main() {
    const auto &&val = get();
}
```



Exercise: Build your experiment library

The above exercise is perfect for building into a set of experiments that are saved in your GitHub gists mentioned in [C++ Is Not Magic](#)



Exercise: Understand how `auto` and `template deduction` relate

Understand the rules for type deduction of templates and how they relate to `auto`.

Read the section in the C++ Programming Language Standard [dcl.spec.auto].

Resources

- [clang-tidy modernize-use-auto¹](#)
- [Almost Always Auto²](#)

¹<https://clang.llvm.org/extra/clang-tidy/checks/modernize-use-auto.html>

²<https://herbsutter.com/2013/08/12/gotw-94-solution-aaa-style-almost-always-auto/>

15. Prefer ranged-for Loop Syntax Over Old Loops

We'll illustrate this point with a series of examples.

int vs std::size_t when looping.

```
for (int i = 0; i < container.size(); ++i) {  
    // oops mismatched types  
}
```

Mismatched containers while looping.

```
for (auto itr = container.begin();  
     itr != container2.end();  
     ++itr) {  
    // oops, most of us have done this at some point  
}
```

Example of ranged-for loop.

```
for(const auto &element : container) {  
    // eliminates both other problems  
}
```



Never mutate the container itself while iterating inside of a ranged-for loop.



Exercise: Modernize Your Loops

You probably have old-style loops in your code.

1. Apply clang-tidy's modernize-loop-convert check.
2. Look for loops that could not be converted.
 - Loops that could not be converted might represent bugs in the code
 - Loops that could not be converted, but do not have bugs, are good candidates for simplification

Resources

- [clang-tidy modernize-loop-convert](https://clang.llvm.org/extra/clang-tidy/checks/modernize-loop-convert.html)¹

¹<https://clang.llvm.org/extra/clang-tidy/checks/modernize-loop-convert.html>

16. Use auto in ranged for loops

Not using auto can make it easier to have silent mistakes in your code.

Accidental conversions

```
for (const int value : container_of_double) {  
    // accidental conversion, possible warning  
}
```

Accidental slicing

```
for (const base value : container_of_derived) {  
    // accidental silent slicing  
}
```

No problem

```
for (const auto &value : container) {  
    // no possible accidental conversion  
}
```

Prefer:

- `const auto` & for non-mutating loops
- `auto` & for mutating loops
- `auto &&` only when you have to work with weird types like `std::vector<bool>`, or if moving elements out of the container



Exercise: Understand `std::map` and ranged for loops

Understand what this code is doing. Is it making a copy? Why and how?

Accidental copy?

```
std::map<std::string, int> get_map();

using element_type = std::pair<std::string, int>;

for (const element_type & : get_map())
{
}
```



Exercise: Enable ranged-loop related warnings

Make sure `-Wrange-loop-construct` is enabled in your code, which is automatically included with `-Wall`.

17. Prefer Algorithms Over Loops

Algorithms communicate meaning and help us apply the “const All The Things” rule. In C++20, we get ranges, which make algorithms more comfortable to use.

It’s possible, taking a functional approach and using algorithms, that we can write C++ that reads like a sentence.

Algorithms end game

```
const auto has_value
    = std::any_of(begin(container), end(container),
                  greater_than(12));
```

Algorithms end game (C++20)

```
const auto has_value
    = std::any_of(container, greater_than(12));
```

Note that in some [rare cases](#)¹, your [static analysis tools](#) might be able to suggest an algorithm to use.



Exercise: Study existing loops

Next time you are reading through a loop in your codebase, cross-reference it with [the C++ `<algorithm>` header](#)² and try to find an algorithm that applies instead.

¹<https://sourceforge.net/p/cppcheck/wiki/ListOfChecks/>

²<https://en.cppreference.com/w/cpp/algorithm>



This book only barely mentions C++20's ranges. Compilers are just now getting support for ranges as of the publication of this book. Ranges can be composed and have full support for `constexpr`.

Resources

- [GoingNative 2013: Sean Parent “C++ Seasoning”](#)³
- [CppCon 2018: Jonathan Boccara “105 Algorithms in Less Than an Hour”](#)⁴
- [C++ Now 2019: Connor Hoekstra “Algorithm Intuition”](#)⁵
- [Code::Dive 2019: Connor Hoekstra “Better Algorithm Intuition”](#)⁶
- [C++ Weekly Ep 187 “C++20’s `constexpr` Algorithms”](#)⁷
- [C++ Weekly Ep 105 “Learning “Modern” C++ 5: Looping And Algorithms”](#)⁸

³<https://channel9.msdn.com/Events/GoingNative/2013/Cpp-Seasoning>

⁴<https://youtu.be/2olsGf6JlIU>

⁵<https://youtu.be/48gV1SNm3WA>

⁶<https://youtu.be/0z-cv3gartw>

⁷<https://youtu.be/9YWzXSr2onY>

⁸<https://youtu.be/A0-x-Djey-Q>

18. Don't Be Afraid of Templates

Templates are the ultimate DRY principle in C++. Templates can be complicated, daunting, and Turing complete, but they don't have to be. Fifteen years ago, it seemed the prevailing attitude is “templates aren't for normal people.”

Fortunately, this is less true today. And we have more tools today, concepts, generic lambdas, etc.

We're going to build up an example over a few chapters. Let's say we want to write a function that can divide any two values.

Divide doubles.

```
double divide(double numerator, double denominator) {  
    return numerator / denominator;  
}
```

But you don't want all of your divisions to be promoted to double.

Divide floats.

```
float divide(float numerator, float denominator) {  
    return numerator / denominator;  
}
```

And of course, you want to handle some kind of integer values.

Divide ints.

```
int divide(int numerator, int denominator) {  
    return numerator / denominator;  
}
```

Templates were designed for just this scenario.

Basic template usage.

```
template<typename T>  
T divide(T numerator, T denominator) {  
    return numerator / denominator;  
}
```

Most examples on the internet use `T`, like I just did. Don't do that. Give your type a meaningful name.

template parameters with actual names.

```
template<typename Arithmetic>  
Arithmetic divide(Arithmetic numerator, Arithmetic denominator) {  
    return numerator / denominator;  
}
```



Exercise: Keep this chapter in mind while moving on to the next chapter and looking at its exercises.

19. Don't Copy and Paste Code

If you find yourself going to select a block of code and copy it: stop!

Take a step back and look at the code again.

- Why are you copying it?
- How similar will the source be to the destination?
- Does it make sense to make a function?
- Remember, [Don't Be Afraid of Templates](#)

I have found that this simple rule has had the most direct influence on my code quality.

If the result of the paste operation was going in the current function, consider using a lambda.

C++14 style lambdas, with generic (aka auto) parameters, give you a simple and easy to use method of creating reusable code that can be shared with different data types while not having to deal with `template` syntax.



Exercise: Try CPD.

There are a few different copy-paste-detectors that look for duplicated code in your codebase.

For this exercise, download the [PMD CPD tool](#)¹ and run it on your codebase.

¹https://pmd.github.io/latest/pmd_userdocs_cpd.html

If you use Arch Linux, this tool can be installed with AUR. The package is pmd; the tool is pmd-cpd.

Can you identify critical parts of your code that have been copied and pasted? What happens if you find a bug in one version? Will you be sure to see all of the locations that also need to be updated?

Resources

- [Copy-Paste Programming](#)²
- [The Last Line Effect](#)³
- [i will not copy-paste code](#)⁴

²<https://www.viva64.com/en/t/0068/>

³<https://www.viva64.com/en/b/0260/>

⁴https://twitter.com/bjorn_fahller/status/1072432257799987200

20. Follow the Rule of 0

No destructor is always better when it's the correct thing to do. Empty destructors can destroy performance:

- They make the type no longer trivial
- Have no functional use
- Can affect inlining of destruction
- Implicitly disable move operations



If you need a destructor because you are doing resource management or defining a base class with virtual functions, you need to follow the [Rule of 5](#).

`std::unique_ptr` can help you apply the Rule of 0 if you provide a custom deleter.



Exercise: Find Rule of 0 Violations in Your Code

Look for code like this (I guarantee you will find it).

Empty meaningless destructor.

```
struct S {  
    // a bunch of other things  
    ~S() {}  
};
```

or worse:

Forward declared empty meaningless destructor.

```
// file.hpp
struct S {
    ~S();
}

// file.cpp
S::~S() {}
```

Are these destructors necessary? Remove them if they are not.

If these destructors exist in types used in many places, you will likely be able to measure smaller binary sizes and better performance by taking this simple action.

Some uses of the plmpl idiom require you to define a destructor. In this case, be sure to follow the [Rule of 5](#).

Resources

- [C++ Reference: The rule of three/five/zero](#)¹
- [C++ Weekly Ep 154: “One Simple Trick for Reducing Code Bloat”](#)²
- [CppCon 2019: Jason Turner “Great C++ is_trivial”](#)³

¹https://en.cppreference.com/w/cpp/language/rule_of_three

²<https://youtu.be/D8eCPl2zit4>

³<https://youtu.be/ZxWjii99yao>

21. If You Must Do Manual Resource Management, Follow the Rule of 5

If you provide a destructor because `std::unique_ptr` doesn't make sense for your use case, you *must* `=delete`, `=default`, or implement the other special member functions.

This rule was initially known as the Rule of 3 and is known as the Rule of 5 after C++11.

The special member functions.

```
struct S {  
    S(); // default constructor  
        // does not affect other special member functions  
  
    // If you define any of the following, you must deal with  
    // all the others.  
    S(const S &);           // copy constructor  
    S(S&&);                 // move constructor  
    S &operator=(const S &); // copy assignment operator  
    S &operator=(S &&);      // move assignment operator  
};
```



`=delete` is a safe way of dealing with the special member functions if you don't know what to do with them!

You should also follow the Rule of 5 when declaring base classes with virtual functions.

Rule of 5 with polymorphic types.

```
struct Base {  
    virtual void do_stuff();  
  
    // because of the virtual function we know this class  
    // is intended for polymorphic use, therefore our  
    // tools will tell us to define a virtual destructor  
    virtual ~Base() = default;  
  
    // and now we need to declare the other special members  
    // a good safe bet is to delete them, because properly and safely  
    // copying or assigning an object via a reference or pointer  
    // to a base class is hard / impossible  
  
    S(S&&) = delete;  
    S(const &S) = delete;  
    S &operator=(const S &) = delete;  
    S &operator=(S &&) = delete;  
};  
  
struct Derived : Base {  
    // We don't need to define any of the special members  
    // here, they are all inherited from `Base`.  
}
```



Instead of `= delete` you can consider making these special members protected.



Exercise: Implement your own `unique_ptr<>` template

It's hard to get it 100% right. Write tests. Understand why the defaulted special member functions don't work.

Bonus points: implement it with C++20's `constexpr` dynamic allocation support.



Exercise: Look for Rule of 5 violations in your code

You are likely not providing consistent lifetime semantics in your existing code when you are defining the special member functions. To assess the impact, you can quickly `= delete;` any missing special member functions and see what breaks.

Resources

- [C++ Reference: The rule of three/five/zero](https://en.cppreference.com/w/cpp/language/rule_of_three)¹

¹https://en.cppreference.com/w/cpp/language/rule_of_three

22. Don't Invoke Undefined Behavior

Ok, there's a lot that's Undefined Behavior (UB), and it's hard to keep track of, so we'll give some examples in the following sections. The critical thing that you need to understand is that UB's existence breaks your entire program.

[[intro.abstract](#)¹]

A conforming implementation executing a well-formed program shall produce the same observable behavior as one of the possible executions of the corresponding instance of the abstract machine with the same program and the same input.

However, if any such execution contains an undefined operation, this document places no requirement on the implementation executing that program with that input (not even with regard to operations preceding the first undefined operation).

Note the sentence “this document places no requirement on the implementation executing that program with that input (not even with regard to operations preceding the first undefined operation)”

If you have UB, the entire program is suspect.



The next several items discuss ways to reduce the risk of undefined behavior in your project.

¹<http://eel.is/c++draft/intro.compliance#intro.abstract-5>



Exercise: Using UBSan, ASan and Warnings

Understanding all of Undefined Behavior is likely impossible. Fortunately, we do have tools that help. Hopefully, you already have your code enabled for UBSan, ASan, and have your warnings enabled. Now is a great time to go back and evaluate what options you have and see if there is anything new you can discover.

Resources

- C++Now 2018: John Regehr “Closing Keynote: Undefined Behavior and Compiler Optimizations”²
- CppCon 2018: Barbara Geller & Ansel Sermersheim “Undefined Behavior is Not an Error”³

²<https://youtu.be/AeEwxtEOgH0>

³https://youtu.be/XEXpwis_deQ

23. Never Test for `this` To Be `nullptr`, It's UB

Invalid check for `this` to be `nullptr`.

```
int Class::member() {  
    if (this == nullptr) {  
        // removed by the compiler, it would be UB  
        // if this were ever null  
        return 42;  
    } else {  
        return 0;  
    }  
}
```

Technically it isn't the check that is Undefined Behavior (UB). But it's impossible for the check ever to fail. If the `this` were to be equal to `nullptr`, you would be in a state of Undefined Behavior.

People used to do this all the time, but it's always been UB. You cannot access an object outside its lifetime. Compilers today will always remove this check.

The only way it's theoretically possible for `this` to be null is when you call a member directly on a null object.



Bad examples lie ahead, do not repeat them.

Bad call of member on nullptr.

```
Type *obj = nullptr;  
obj->do_thing(); // never do this
```

Even in the (technically OK, but never do this) scenario of calling `delete this`.

Bad example of `delete this`.

```
struct S {  
    std::string data;  
  
    void delete_yourself() {  
        // do things  
        delete this; // technically OK  
  
        if (this) {  
            // this block will always be executed, nothing changed  
            // our view of `this`  
        }  
  
        // never do this  
        data.size(); // UB, data's lifetime has ended  
    }  
};
```

There is no scenario in which a check for `if (this)` will return false on a modern compiler.



Exercise: Do you check for `this to be nullptr`?

A check for `nullptr` can hide as a check for `NULL` or a check against `0`. A check for `this to be NULL` is likely to only exist in very old code bases. Make sure you have your warnings enabled, then look for these cases.

It's probably interesting in general to search for `this ==` in your codebase and see what weird things are there.

Resources

- [Porting to GCC-6 Optimizations remove null pointer checks for `this`](https://www.gnu.org/software/gcc/gcc-6/porting_to.html#this-cannot-be-null)¹

¹https://www.gnu.org/software/gcc/gcc-6/porting_to.html#this-cannot-be-null

24. Never Test for A Reference To Be `nullptr`, It's UB

Tests for null references are removed

```
int get_value(int &thing) {  
    if (&thing == nullptr) {  
        // removed by compiler  
        return 42;  
    } else {  
        return thing;  
    }  
}
```

It's UB to make a null reference, don't try it. Always assume a reference refers to a valid object. Use this fact to your advantage when [designing API's](#).



Exercise: Check for checking the address of an object

There are many valid use cases for `&thing ==` to check for a specific address of an object, but there are also many ways this check can be wrong.

Search through your code for statements that check an object's memory address and understand what they are doing and how (or if) they work.



What other ways might the address of an object be checked besides `==`?

This exercise gives you some great experience working with various searching / grepping tools and playing with regex.

Resources

- [-Wtautological-undefined-compare](https://clang.llvm.org/docs/DiagnosticsReference.html#wtautological-undefined-compare)¹

¹<https://clang.llvm.org/docs/DiagnosticsReference.html#wtautological-undefined-compare>

25. Avoid default In switch Statements

This is an issue that is best described with a series of examples. Starting from this one:

switch with warnings

```
enum class Values {
    val1,
    val2
};

std::string_view get_name(Values value) {
    switch (value) {
        case val1: return "val1";
        case val2: return "val2";
    }
}
```

If you have enabled all of your warnings, then you will likely get a “not all code paths return a value” warning here. Which is technically correct. We could call `get_name(static_cast<Values>(15))` and not violate any part of C++ [dcl.enum/5] except for the Undefined Behavior of not returning a value from a function.

You'll be tempted to fix this code like this:

switch with default to avoid warnings

```
enum class Values {
    val1,
    val2
};

std::string_view get_name(Values value) {
    switch (value) {
        case Values::val1: return "val1";
        case Values::val2: return "val2";
        default: return "unknown";
    }
}
```

But this introduces a new problem

Unhandled case

```
enum class Values {
    val1,
    val2,
    val3 // added a new value
};

std::string_view get_name(Values value) {
    switch (value) {
        case Values::val1: return "val1";
        case Values::val2: return "val2";
        default: return "unknown";
    }
    // no compiler diagnostic that `val3` is unhandled
}
```

Instead, prefer code like this:

Preferred version

```
enum class Values {
    val1,
    val2,
    val3 // added a new value
};

std::string_view get_name(Values value) {
    switch (value) {
        case Values::val1: return "val1";
        case Values::val2: return "val2";
    } // unhandled enum value warning now

    return "unknown";
}
```



You shouldn't ever get an "unreachable code" warning in the above example because the [range of valid values](#) is nearly always larger than the values you have defined.



Some modern tools can detect these uses of default for you.



Exercise: Look for default:.

What do you find in your code base? Did enabling warnings in previous exercises find uses of default: for you already?

Resources

- CppCon 2018: Jason Turner “Applied Best Practices”¹
- `-Wswitch-enum`²
- `-Wswitch`³

¹<https://youtu.be/DH0IsEd0eDE>

²<https://clang.llvm.org/docs/DiagnosticsReference.html#wswitch-enum>

³<https://clang.llvm.org/docs/DiagnosticsReference.html#wswitch>

26. Prefer Scoped enums

C++11 introduced scoped enumerations, intended to solve many of the common problems with enum inherited from C.

C++98 enums

```
enum Choices {
    option1 // value in the global scope
};

enum OtherChoices {
    option2
};

int main() {
    int val = option1;
    val = option2; // no warning
}
```

- enum Choices;
- enum OtherChoices;

These two can easily get mixed up, and they each introduce identifiers in the global namespace.

- enum class Choices;
- enum class OtherChoices;

The values in these enumerations are scoped and more strongly typed.

C++11 scoped enumeration.

```
enum class Choices {
    option1
};

enum class OtherChoices {
    option2
};

int main() {
    int val = option1; // cannot compile, need scope
    int val2 = Choices::option1; // cannot compile, wrong type
    Choices val = Choices::option1; // compiles
    val = OtherChoices::option2; // cannot compile, wrong type
}
```

These `enum class` versions cannot get mixed up without much effort, and their identifiers are now scoped, not global.

`enum struct` and `enum class` are equivalent. Logically `enum struct` makes more sense since they are public names. Which do you prefer?



Exercise: `enum struct` OR `enum class`

Decide if you prefer `enum struct` or `enum class` and develop a well-reasoned answer as to why.



Exercise: clang-tidy modernize

Clang-tidy's modernizer can add `class` to your `enum` declarations. Try putting it to use.



clang-tidy's scoped enumeration modernizer will probably find many bugs in your code!

Resources

- [CppCon 2018: Victor Ciura “Better Tools in Your Clang Toolbox”](#)¹ (Discusses bugs found by moving to `enum class`)
- [cppreference.com Enumeration Declaration](#)²

¹https://youtu.be/4X_fZkl7kkU

²<https://en.cppreference.com/w/cpp/language/enum>

27. Prefer `if constexpr` over SFINAE

SFINAE is kind-of write-only code. `if constexpr` doesn't have the same flexibility, but use it when you can.

Let's take our divide example last seen in [Prefer `auto` in Many Cases](#):

C++14 divides template.

```
template<typename Numerator, typename Denominator>
auto divide(Numerator numerator, Denominator denominator)
{
    return numerator / denominator;
}
```

We now want to add different behavior if we are doing integral division. Before C++17, we would have used SFINAE (“Substitution Failure Is Not An Error”). Essentially this means that if a function fails to compile, then it is removed from overload resolution.

SFINAE divide function.

```
#include <stdexcept>
#include <type_traits>
#include <utility>

template <typename Numerator, typename Denominator,
          std::enable_if_t<std::is_integral_v<Numerator> &&
                          std::is_integral_v<Denominator>,
                          int> = 0>
auto divide(Numerator numerator, Denominator denominator) {
    // is integer division
    if (denominator == 0) {
        throw std::runtime_error("divide by 0!");
    }
    return numerator / denominator;
}

template <typename Numerator, typename Denominator,
          std::enable_if_t<std::is_floating_point_v<Numerator> ||
                          std::is_floating_point_v<Denominator>,
                          int> = 0>
auto divide(Numerator numerator, Denominator denominator) {
    // is floating point division
    return numerator / denominator;
}
```

The `if constexpr` construct in C++17 can simplify this code:

`if constexpr` *option for compile time behavior change.*

```
#include <stdexcept>
#include <type_traits>
#include <utility>

template <typename Numerator, typename Denominator>
auto divide(Numerator numerator, Denominator denominator) {
    if constexpr (std::is_integral_v<Numerator> &&
                  std::is_integral_v<Denominator>) {
        // is integral division
        if (denominator == 0) {
            throw std::runtime_error("divide by 0!");
        }
    }

    return numerator / denominator;
}
```

Note that the code inside the `if constexpr` block must still be syntactically correct. `if constexpr` is not the same as a `#define`.

Resources

- C++ Weekly Special Edition: Using C++17's `constexpr if`¹
- C++ Weekly Ep 122: `constexpr` with `optional` and `variant`²
- CppCon 2017: Jason Turner “Practical C++17”³
- C++17 In Tony Tables: `constexpr if`⁴

¹https://youtu.be/_Ny6Qbm_uMI

²https://youtu.be/2eCV_udkP_o

³<https://youtu.be/nnY4e4faNp0>

⁴https://github.com/tvaneerd/cpp17_in_TTs/blob/master/if_constexpr.md

28. Constrain Your Template Parameters With Concepts (C++20)

Concepts will result in better error messages (eventually) and better compile times than SFINAE. Besides much more readable code than SFINAE.

If we continue to build on our divide example, we can take this `if constexpr` version from the [Prefer if constexpr over SFINAE](#) chapter.

if constexpr version of divide function.

```
#include <stdexcept>
#include <type_traits>
#include <utility>

template <typename Numerator, typename Denominator>
auto divide(Numerator numerator, Denominator denominator) {
    if constexpr (std::is_integral_v<Numerator> &&
                  std::is_integral_v<Denominator>) {
        // is integral division
        if (denominator == 0) {
            throw std::runtime_error("divide by 0!");
        }
    }

    return numerator / denominator;
}
```

And we can split it back out as two different functions using concepts.

Concepts can be used in several different contexts. This version uses a simple `requires` clause after the function declaration.

Concepts in requires clause.

```
#include <stdexcept>
#include <type_traits>
#include <utility>

// overload resolution will pick the most specific version
template <typename Numerator, typename Denominator>
auto divide(Numerator numerator, Denominator denominator) requires
    (std::is_integral_v<Numerator>
     && std::is_integral_v<Denominator>) {
    // is integral division
    if (denominator == 0) {
        throw std::runtime_error("divide by 0!");
    }
    return numerator / denominator;
}

template <typename Numerator, typename Denominator>
auto divide(Numerator numerator, Denominator denominator) {
    return numerator / denominator;
}
```

This version uses concepts as function parameters. C++20 even has an “auto concept,” which is an implicit template function.

Terse concepts requirement syntax.

```
#include <stdexcept>
#include <concepts>

auto divide(std::integral auto numerator,
            std::integral auto denominator) {
    // is integer division
    if (denominator == 0) {
        throw std::runtime_error("divide by 0!");
    }
}
```

```
    return numerator / denominator;
}

auto divide(auto numerator, auto denominator) {
    // is floating point division
    return numerator / denominator;
}
```



Concepts can define complex requirements, including expected members. This section only barely touches on the possibilities.



Exercise: Understand what concepts are provided with C++20.

As usual, `cppreference` helps [by providing a list of concepts](https://en.cppreference.com/w/cpp/concepts)¹.



Exercise: Create your own concept.

Does this example give you some idea for an example of a concept that you would want, but isn't provided by `<concepts>`?

Look at the implementation of the very simple `std::integral` concept on [cppreference](https://en.cppreference.com/w/cpp/concepts/integral)² and see if it inspires you.

¹<https://en.cppreference.com/w/cpp/concepts>

²<https://en.cppreference.com/w/cpp/concepts/integral>

Resources

- [C++ Weekly Ep 194: From SFINAE To Concepts With C++20](#)³
- [C++ Weekly Ep 196: What is `requires` `requires`](#)⁴

³<https://youtu.be/dR64GQb4AGo>

⁴https://youtu.be/tc0hVIOJk_U

29. De-template-ize Your Generic Code

Move things outside of your templates when you can. Use other functions. Use base classes. The compiler is still free to inline them or leave them out of line.

De-template-ization will improve compile times and reduce binary sizes. Both are helpful. It also eliminates the thing that people think of as “template code bloat” (which IMO [doesn't exist](#)¹) (article formatting got broken at some point, sorry).

A new lambda for each function template instantiation.

```
template<typename T>
void do_things()
{
    // this lambda must be generated for each
    // template instantiation
    auto lambda = [](){ /* some lambda that doesn't capture */ };
    auto value = lambda();
}
```

Compared to:

¹https://articles.emptycrate.com/2008/05/06/nobody_understands_c_part_5_template_code_bloat.html

Shared logic between template instantiations.

```
auto some_function() { /* do things*/ }

template<typename T>
void do_things()
{
    auto value = some_function();
}
```

Now only one version of the inner logic is compiled, and it's up to the compiler to decide if they should be inlined.

Similar techniques apply to base classes and templated derived classes.



Exercise: Bloaty McBloatface and -ftime-trace.

We're getting more and more tools available to look for bloat in our binaries and analyze compile times. Look into these tools and other tools available on your platform.

Run them against your binary and see what you find.

When using clang's `-ftime-trace`, also look into ClangBuildAnalyzer.

Resources

- [Templight](https://github.com/mikael-s-persson/templight)²
- [C++ Weekly Ep 89: "Overusing Lambdas"](https://youtu.be/OmKMNQFx_8Y)³

²<https://github.com/mikael-s-persson/templight>

³https://youtu.be/OmKMNQFx_8Y

- [C++ Weekly Christmas Class 2019 - Chapter 3](#)⁴ (This is the first episode of chapter 3, and it introduces the question of how and why two different [options differ](#)⁵. The next several episodes in that playlist give some background, and the start of chapter 4 gives the answers. It is very much related to template bloat questions.)
- Effective C++ (3rd Edition) Item 44 - Factor parameter-independent code out of templates

⁴https://www.youtube.com/watch?v=VEqOOKU8RjQ&list=PLs3KjaCtOwSY_Awyliwm-fRjEOa-SRbs-&index=16

⁵<https://godbolt.org/z/b4znvK>

30. Use Lippincott Functions

Same arguments as [de-template-izing](#) your code: This is a do-not-repeat-yourself principle for exception handling routines.

If you have many different exception types to handle, you might have code that looks like this:

Duplicated exception handling.

```
void use_thing() {
    try {
        do_thing();
    } catch (const std::runtime_error &) {
        // handle it
    } catch (const std::exception &) {
        // handle it
    }
}
```

```
void use_other_thing() {
    try {
        do_other_thing();
    } catch (const std::runtime_error &) {
        // handle it
    } catch (const std::exception &) {
        // handle it
    }
}
```

A Lippincott function (named after Lisa Lippincott) provides a centralized exception handling routine.

Lippincott de-duplicated exception handling.

```
void handle_exception() {
    try {
        throw; // re-throw exception already in flight
    } catch (const std::runtime_error &) {
    } catch (const std::exception &) { }
}

void use_thing() {
    try {
        do_thing();
    } catch (...) {
        handle_exception();
    }
}

void use_other_thing() {
    try {
        do_other_thing();
    } catch (...) {
        handle_exception();
    }
}
```

This technique is not new - it has been available since the pre-C++98 days.



Exercise: Do You Use Exceptions?

If your project uses exceptions, there's probably some ground for simplifying and centralizing your error handling routines. If it does not use exceptions, then you likely have other types of error handling routines that are duplicated. Can these be simplified?

Resources

- [C++ Secrets: Using a Lippincott Function for Centralized Exception Handling](#)¹
- [C++ Weekly Ep 91: Using Lippincott Functions](#)²

¹<https://cppsecrets.blogspot.com/2013/12/using-lippincott-function-for.html>

²<https://youtu.be/-amJL3AyADI>

31. Be Afraid of Global State

Reasoning about global state is hard.

Any non-const static value, or `std::shared_ptr<>` could potentially be global state. It is never known who might update the value or if it is thread-safe to do so.

Global state can result in subtle and difficult to trace bugs where one function changes global state, and another function either relies on that change or is adversely affected by it.



Exercise: Global State, What's Left?

If you've done the other exercises, you've already made all of your static variables const. This is great! You've possibly even made some of them `constexpr`, which is even better!

But you probably have global state still lurking. Do you have a global singleton logger? Could the logger be accidentally sharing state between the modules of your system?

What about other singletons? Can they be eliminated? Do they have threading initialization issues (what happens if two threads try to access one of the objects for the first time at the same time)?

Resources

- [Retiring the Singleton Pattern - Peter Muldoon - Meeting C++ 2019](#)¹

¹<https://youtu.be/f46jmm7r8Yg>

32. Make your interfaces hard to use wrong.

Your interface is your first line of defense. If you provide an interface that is easy to use wrong, your users *will* use it wrong.

If you provide an interface that's hard to use wrong, your users have to work harder to use it wrong. But this is still C++; they will always find a way.

Interfaces hard to use wrong will sometimes result in more verbose code where we would maybe like more terse code. You have to choose what is most important. Correct code or short code?

This is a high-level concept; specific ideas will follow.

Resources

* [The Little Manual of API Design](#)¹

¹<https://people.mpi-inf.mpg.de/~jblanche/api-design.pdf>

33. Consider If Using the API Wrong Invokes Undefined Behavior

Do you accept a raw pointer? Is it an optional parameter? What happens if `nullptr` is passed to your function?

What happens if a value out of the expected range is passed to your function?

Some developers make the distinction between “internal” and “external” APIs. They allow unsafe APIs for internal use only.



Is there any guarantee that an external user will never invoke the “internal” API?

Is there any guarantee that your internal users will never misuse the API?



Exercise: Investigate Checked Types

The C++ Guideline Support Library (GSL) has a `not_null` pointer type that guarantees, because of zero cost abstractions, that the pointer passed is never `nullptr`. Would that work for your APIs that currently pass raw pointers (assuming that rearchitecting the API is not an option)?

`std::string_view` (C++17) and `std::span` (C++20) are great alternatives to pointer / length pairs passed to functions.

Resources

- [boost::safe_numerics](https://github.com/boostorg/safe_numerics)¹

¹https://github.com/boostorg/safe_numerics

34. Use `[[nodiscard]]` Liberally

`[[nodiscard]]` (C++17) is a C++ attribute that tells the compiler to warn if a return value is ignored. It can be used on functions:

`[[nodiscard]]` *example usage.*

```
[[nodiscard]] int get_value();

int main()
{
    // warning, [[nodiscard]] value ignored
    get_value();
}
```

And on types:

`[[nodiscard]]` *on types.*

```
struct [[nodiscard]] ErrorCode{};

ErrorCode get_value();

int main()
{
    // warning, [[nodiscard]] value ignored
    get_value();
}
```

C++20 adds the ability to provide a description.

C++20's `[[nodiscard]]` *with description*.

```
[[nodiscard("Ignoring this result leaks resources")]]
```

Our divide example is a straightforward application of `[[nodiscard]]`.

`[[nodiscard]]` *applied to divide function*.

```
#include <stdexcept>
#include <concepts>

[[nodiscard]] auto divide(std::integral auto numerator,
                          std::integral auto denominator) {
    // is integer division
    if (denominator == 0) {
        throw std::runtime_error("divide by 0!");
    }
    return numerator / denominator;
}

[[nodiscard]] auto divide(auto numerator, auto denominator) {
    // is floating point division
    return numerator / denominator;
}
```



Exercise: Determine a set of rules for using `[[nodiscard]]`

Read the Reddit discussion “[An Argument Pro Liberal Use Of `nodiscard`](https://www.reddit.com/r/cpp/comments/9us7f3/an_argument_pro_liberal_use_of_nodiscard/)”¹. Consider your types and functions. Which values should be `[[nodiscard]]`?

Should it be a compiler error or warning to call these functions and ignore the result?

¹https://www.reddit.com/r/cpp/comments/9us7f3/an_argument_pro_liberal_use_of_nodiscard/

- `vector.size()`
- `vector.empty()`
- `vector.insert()`

Resources

- “An Argument Pro Liberal Use Of `nodiscard`”²
- C++ Weekly Ep 30: C++17’s `[[nodiscard]]` Attribute³
- C++ Weekly Ep 199: C++20’s `[[nodiscard]]` Constructors And Their Uses⁴

²https://www.reddit.com/r/cpp/comments/9us7f3/an_argument_pro_liberal_use_of_nodiscard/

³https://youtu.be/_5PF3GQLKc

⁴https://youtu.be/E_ROB_xUQQQ

35. Use Stronger Types

Consider the API for POSIX socket:

POSIX socket API.

```
socket(int, int, int);
```

The parameters (in some order) represent:

- type
- protocol
- domain

This design is problematic, but there are less obvious ones lurking in our code.

Poorly defined constructor.

```
Rectangle(int, int, int, int);
```

This function could be (x, y, width, height), or (x1, y1, x2, y2). Less likely, but still possible, is (width, height, x, y).

What do you think about an API that looks like this?

Strongly typed constructor.

```
Rectangle(Position, Size);
```

In many cases, it only takes a little effort to make more strongly typed APIs.

Stronger typed definitions.

```
struct Position {  
    int x;  
    int y;  
};  
  
struct Size {  
    int width;  
    int height;  
};  
  
struct Rectangle {  
    Position position;  
    Size size;  
};
```

Which can then lead to other, logically composable statements with operator overloads such as:

Coupled type operator overload.

```
// Return a new rectangle that has been  
// moved by the offset amount passed in  
Rectangle operator+(Rectangle, Position);
```



It's possible that making structs can *increase* performance in some cases [C++ Weekly Ep 119, Negative Cost Structs](#)¹.

Avoid Boolean Arguments

This chapter's pre-release reader pointed out that Steve Maguire says, "Make code intelligible at the point of call. Avoid Boolean arguments," in Chapter 5 of his book *Writing Solid Code*.

¹<https://youtu.be/FwsO12x8nyM>

In C++11, `enum class` gives you an easy way to add stronger typing, avoid boolean parameters, and make your API harder to use wrong.

Consider:

Non-obvious order of parameters.

```
struct Widget {  
    // this constructor is easy to use wrong, we  
    // can easily transpose the parameters  
    Widget(bool visible, bool resizable);  
}
```

Compared to:

Stronger typing with scoped enumerations.

```
struct Widget {  
    enum struct Visible { True, False };  
    enum struct Resizable { True, False };  
  
    // still possible to use this wrong, but MUCH harder  
    Widget(Visible visible, Resizable resizable);  
}
```



Identify the problematic APIs in your existing code.

What function call do you regularly get out of order? How can it be fixed?



Exercise: Research strong typedef libraries for C++.

There are existing libraries that simplify some of the boilerplate code for you when making a strongly typed `int`. Jonathan Muller, Bjorn Fahlber, and Peter Sommerlad have each written one, and others are available.



Exercise: Consider `=delete`ing problematic conversions.

Simple function declaration.

```
double high_precision_thing(double);
```

What if calling the above with a `float` is likely to be a bug?

Deleting a problematic accidental promotion from `float` to `double`.

```
double high_precision_thing(double);  
double high_precision_thing(float) = delete;
```

Any function or overload can be `=deleted` in C++11.

Resources

- C++ Weekly Ep 107: “The Power of `=delete`”²
- Adi Shavit and Björn Fahlber “The Curiously Recurring Pattern of Coupled Types”³
- Research “Affine space types.”
- C++Now 2017: Jonathan Müller “Type-safe Programming”⁴

²<https://youtu.be/aAvjUU0m6AU>

³<https://youtu.be/msi4WNQZyWs>

⁴<https://youtu.be/iihlo9A2Ezw>

36. Don't return raw pointers

Returning a raw pointer makes the reader of the code and user of the library think too hard about ownership semantics. Prefer a reference, smart pointer, non owning pointer wrapper, or consider an optional reference.

Function returning a raw pointer.

```
int *get_value();
```

Who owns this return value? Do I? Is it my job to delete it when I'm done with it?

Or even worse, what if the memory was allocated by `malloc` and I need to call `free` instead?

Is it a single `int` or an array of `int`?

This code has far too many questions, and not even `[[nodiscard]]` can help us.



Exercise: You know the drill

By now, you've done enough of these API related exercises to know what to do. Go and look for these in your code! See if there's a better way! Can you return a value, reference, or `std::unique_ptr` instead?

37. Prefer Stack Over Heap

Stack objects (locally scoped objects that are not dynamically allocated) are much more optimizer friendly, cache-friendly, and may be entirely eliminated by the optimizer. As Björn Fahlner [has said](#), “assume any pointer indirection is a cache miss.”

In the most simple terms:

OK idea, uses stack and can be optimized.

```
std::string make_string() { return "Hello World"; }
```

Bad idea, uses the heap.

```
std::unique_ptr<std::string> make_string() {  
    return std::make_unique<std::string>("Hello World");  
}
```

OK idea.

```
void use_string() {  
    // This string lives on the stack  
    std::string value("Hello World");  
}
```

Really bad idea, uses the heap and leaks memory.

```
void use_string() {  
    // The string lives on the heap  
    std::string *value = new std::string("Hello World");  
}
```



Remember, `std::string` itself might allocate internally, and use the heap. If no heap usage at all is your goal, you will need to take other measures. The goal is no *unnecessary* heap allocations.

Generally speaking, objects created with `new` expressions (or via `make_unique` or `make_shared`) are heap objects, and have *Dynamic Storage Duration*. Objects created in a local scope are stack objects and have *Automatic Storage Duration*.



Exercise: Look for heap usage

Sometimes developers with C and Java backgrounds have a hard time with this. For Java, it's because `new` is required to create objects. For C, it is because the C compiler cannot perform the same kinds of optimizations that the C++ compiler can because of differences in the language.

So some of this unnecessary heap usage may have ended up in your current code.



Exercise: Run a heap profiler

There are several heap profiling tools, and there may even be one built into your IDE. Examine your heap usage and look for potential abuses of the heap in your project.

Resources

- [Code::Dive 2018: Björn Fahlner “What Do You Mean By Cache Friendly?”](#)¹

¹<https://youtu.be/Fzbotzi1gYs>

38. No More new!

You're already [avoiding the heap](#) and using smart pointers for resource management, right?!

Take this to the next level and be sure to use `std::make_unique<>()`¹ (C++14) in the rare cases that you need the heap.

In the very rare cases you need shared ownership, use `std::make_shared<>()`² (C++11).



Exercise: Do you use Qt or some other widget library?

Have you ever thought about writing your own `make_qobject` helper? Give it the semantics you need and be sure to use `[[nodiscard]]`.

In any case, you can limit your use of `new` to a few core library helper functions.



Exercise: Use clang-tidy modernize fixes.

With clang-tidy, you can automatically convert `new` statements into `make_unique<>` and `make_shared<>` calls. Be sure to use `-fix` to apply the change after it's been discovered.

¹https://en.cppreference.com/w/cpp/memory/unique_ptr/make_unique

²https://en.cppreference.com/w/cpp/memory/shared_ptr/make_shared

Resources

- [clang-tidy modernize-make-shared](https://clang.llvm.org/extra/clang-tidy/checks/modernize-make-shared.html)³
- [clang-tidy modernize-make-unique](https://clang.llvm.org/extra/clang-tidy/checks/modernize-make-unique.html)⁴

³<https://clang.llvm.org/extra/clang-tidy/checks/modernize-make-shared.html>

⁴<https://clang.llvm.org/extra/clang-tidy/checks/modernize-make-unique.html>

39. Know Your Containers

Prefer your containers in this order:

- `std::array<>`
- `std::vector<>`

`std::array<>`

A fixed-size stack-based contiguous container. The data size must be known at compile-time, and you must have enough stack space to hold the data. This container helps us [prefer stack over heap](#). Known location and contiguousness results in `std::array<>` becoming a “negative cost abstraction.” The compiler can perform an extra set of optimizations because it knows the data’s size and location.

`std::vector<>`

A dynamically-sized heap-based contiguous container. While the compiler does not know where the data will ultimately reside, it does know that the elements are laid out adjacent to each other in RAM. Contiguousness gives the compiler more optimization opportunities and is more [cache-friendly](#).

Almost anything else needs a comment and justification for why. A flat map with linear search is likely better than an `std::map` for small containers.

But don’t be too enthusiastic about this. If you need key lookup, use `std::map` and evaluate if it has the performance and characteristics you want.



Exercise: Replace vector With array

Look for fixed-size vectors and replace them with array where possible. With C++17's Class Template Argument Deduction, this can be easier.

`const` `std::vector` with fixed-size data.

```
const std::vector<int> data{n+1, n+2, n+3, n+4};
```

can become

`const` `std::array` for fixed-size data.

```
const std::array<int, 4> data{n+1, n+2, n+3, n+4}; // C++11  
const std::array data{n+1, n+2, n+3, n+4};         // C++17
```

You already made these `const`, now go back to `constexpr` them if you can.

Resources

- Bjarne Stroustrup “Are lists evil?”¹

¹https://www.stroustrup.com/bs_faq.html#list

40. Avoid `std::bind` and `std::function`

While compilers continue to improve and the optimizers work around these types' complexity, it's still very possible for either to add considerable compile-time and runtime overhead.

C++14 lambdas, with generalized capture expressions, are capable of the same things that `std::bind` is capable of.

std::bind to change parameter order

```
#include <functional>
```

```
double divide(double numerator, double denominator) {  
    return numerator / denominator;  
}
```

```
auto inverted_divide = std::bind(divide,  
                                std::placeholders::_2,  
                                std::placeholders::_1);
```

Lambda to change parameter order

```
#include <functional>

double divide(double numerator, double denominator) {
    return numerator / denominator;
}

auto inverted_divide = [](const auto numerator,
                          const auto denominator) {
    return divide(denominator/numerator)
};
```



Exercise: Compare the possibilities.

Take these options in Compiler Explorer. How do the compile times and resulting assembly look?

`std::function` *and* `std::bind`

```
#include <functional>

template<typename Func>
std::function<int (int)> bind_3(Func func)
{
    return std::bind(func, std::placeholders::_1, 3);
}

int main(int argc, const char *[])
{
    return bind_3(std::plus<>{})(argc);
}
```

std::bind only, for bonus points, what type is returned from the function `bind_3`?

```
#include <functional>

template<typename Func>
auto bind_3(Func func)
{
    return std::bind(func, std::placeholders::_1, 3);
}

int main(int argc, const char *[])
{
    return bind_3(std::plus<>{})(argc);
}
```

Only lambdas, no std library wrappers.

```
#include <functional>

template<typename Func>
auto bind_3(Func func)
{
    return [func](const int value){ return func(value, 3); };
}

int main(int argc, const char *[])
{
    return bind_3(std::plus<>{})(argc);
}
```

Resources

- CppCon 2015: Stephan T. Lavavej “<functional>: What’s New, And Proper Usage”¹
- C++ Weekly Ep 16: “Avoiding `std::bind`”²

¹<https://youtu.be/zt7ThwVfap0>

²<https://youtu.be/ZlHi8txU4aQ>

41. Skip C++11

If you're currently looking to move to "modern" C++, finally, please skip C++11. C++14 fixes several holes in C++11.

Language Features

- C++11's version of `constexpr` implies `const` for member functions, this is changed in C++14
- C++11 is missing auto return type deduction for regular functions (lambdas have it)
- C++11 does not have auto or variadic lambda parameters
- C++14 adds `[[deprecated]]` attribute
- C++14 adds ' digit separator, example: `1'000'000`
- `constexpr` functions can be more than just a single return statement in C++14

Library Features

- `std::make_unique` was added in C++14, which enables the "no raw new" standard
- C++11 doesn't have `std::exchange`
- C++14 adds some `constexpr` support for `std::array`
- `cbegin`, `cend`, `crbegin` and `crend` free functions added for consistency with `begin` and `end` free functions and member functions added to standard containers in C++11.



Exercise: Can you use C++14 today?

Even if your project is currently stuck on an older compiler released before C++14, it is highly likely that you can use C++14 features if you enable `-std=c++1y` or `-std=c++14` mode.

Compare the [C++14 language feature chart](https://en.cppreference.com/w/cpp/compiler_support/14)¹ from [cppreference.com](https://en.cppreference.com/w/cpp/compiler_support/14) to the compiler you currently require for your project. How many features could you be taking advantage of today?

As of GCC 5, all of C++14 is supported, but as early as 4.9 provided many C++14 features.



Exercise: Can you go beyond C++14?

Ask if it's possible to upgrade your current compiler requirements. With very rare exceptions, each new compiler version brings:

- Better performance
- Fewer bugs
- Better warnings
- Better standards conformance



A few features were removed from C++17 such as `std::auto_ptr`, `std::unary_function` and `std::binary_function`. You may run into these issues when moving your project to C++17 mode. Most uses of `std::unary_function` and `std::binary_function` can be removed with no change to the rest of the code.

¹https://en.cppreference.com/w/cpp/compiler_support/14

Resources

- C++ Weekly Ep 173: The Important Parts of C++98 in 13 Minutes²
- C++ Weekly Ep 176: The Important Parts of C++11 in 12 Minutes³
- C++ Weekly Ep 178: The Important Parts of C++14 in 9 Minutes⁴
- C++ Weekly Ep 190: The Important Parts of C++17 in 10 Minutes⁵

²https://youtu.be/78Y_LRZPVRg

³<https://youtu.be/D5n6xMUKU3A>

⁴<https://youtu.be/mXxNvaEdNHI>

⁵<https://youtu.be/QpFjOlzg1r4>

42. Don't Use `initializer_list` For Non-Trivial Types

“Initializer List” is an overloaded term in C++. “Initializer Lists” are used to directly initialize values. `initializer_list` is used to pass a list of values to a function or constructor.



Exercise: Understand the overhead `initializer_list` can bring

Use Andreas Fertig's awesome cppinsights.io¹ to understand what these two examples do

`initializer_list` constructor with `shared_ptr`.

```
#include <vector>
#include <memory>

std::vector<std::shared_ptr<int>>> vec{
    std::make_shared<int>(40), std::make_shared<int>(2)
};
```

¹<http://cppinsights.io>

std::array construction with shared_ptr.

```
#include <array>
#include <memory>

std::array<std::shared_ptr<int>, 2> data{
    std::make_shared<int>(40), std::make_shared<int>(2)
};
```

And explain the difference. If you can do this, you understand more than most C++ developers.



Exercise: Understand why this doesn't compile

initializer_list construction with unique_ptr.

```
#include <vector>
#include <memory>

std::vector<std::unique_ptr<int>> data{
    std::make_unique<int>(40), std::make_unique<int>(2)
};
```

Resources

- C++Now 2018: Jason Turner “Initializer Lists Are Broken, Let's Fix Them”² (deep dive into the issues around these topics)
- C++ Insights³

²<https://youtu.be/sSImmZMFsXQ>

³<https://cppinsights.io/>

43. Use the Tools: Build Generators

- CMake¹
- Meson²
- Bazel³
- Others⁴

Raw make files or Visual Studio project files make each of the things listed above too tricky to implement. Use a build tool to help you with maintaining portability across platforms and compilers.

Treat your build scripts like any other code. They have their own set of best practices, and it's just as easy to write unmaintainable build scripts as it is to write unmaintainable C++.

Build generators also help abstract and simplify your continuous build environment with tools like `cmake --build`, which does the correct thing regardless of the platform in use.

¹<https://cmake.org>

²<https://mesonbuild.com/>

³<https://bazel.build/>

⁴https://github.com/lefticus/cppbestpractices/blob/master/02-Use_the_Tools_Available.md



Exercise: Investigate your build system.

- Does your project currently use a build generator?
- How old are your build scripts?

See if there are current best practices you need to apply. Are there tidy-like or formatting tools you can run on your build scripts?

Read back over the previous best practices from this book and see how they apply to your build scripts.

- Are you repeating yourself?
- Are there higher-level abstractions available?



Recent versions of CMake have added tools like `--profiling-output` to help you see where the generator is spending its time.

Resources

- [Professional CMake: A Practical Guide](https://crascit.com/professional-cmake/)⁵
- [cmake-tidy](https://github.com/MaciejPatro/cmake-tidy)⁶
- [C++Now 2017: Daniel Pfeiffer “Effective CMake”](https://youtu.be/bsXLMQ6WgIk)⁷
- [CppCon 2017: Mathieu Ropert “Using Modern CMake Patterns to Enforce a Good Modular Design”](https://youtu.be/eC9-iRN2b04)⁸

⁵<https://crascit.com/professional-cmake/>

⁶<https://github.com/MaciejPatro/cmake-tidy>

⁷<https://youtu.be/bsXLMQ6WgIk>

⁸<https://youtu.be/eC9-iRN2b04>

- CppCon 2018: Jussi Pakkanen “Compiling Multi-Million Line C++ Code Bases Effortlessly with the Meson Build System”⁹
- BazelCon 2019¹⁰
- CppCon 2019: Mathieu Ropert “Cato the Elder”¹¹ - Short rant about build script quality
- C++ Weekly Ep 218 - The Ultimate CMake / C++ Quick Start¹²
- Twitter Discussion on CMake Resources¹³

⁹<https://youtu.be/SCZLnopmYBM>

¹⁰<https://www.youtube.com/playlist?list=PLxNYxgaZ8Rsf-7g43Z8LyXct9ax6egdSj>

¹¹<https://youtu.be/D07iF4Lp4kM>

¹²<https://youtu.be/YbgH7yat-Jo>

¹³<https://twitter.com/Cor3ntin/status/1310990444915032067>

44. Use the Tools: Package Managers

Recent years have seen an explosion of interest in package managers for C++. These two have become the most popular:

- [Vcpkg](https://github.com/Microsoft/vcpkg)¹
- [Conan](https://conan.io/)²

There is a definite advantage to using a package manager. Package managers help with portability and reducing maintenance load on developers.



Exercise: What are your dependencies?

Take time to inventory your project's dependencies. Compare your dependencies with what is available with the package managers above. Does any one package manager have all of your dependencies? How out of date are your current packages? What security fixes are you currently missing?

¹<https://github.com/Microsoft/vcpkg>

²<https://conan.io/>

45. Improving Build Time

A few practical considerations for making build time less painful

- [De-template-ize your code where possible](#)
- Use forward declarations where it makes sense to
- Enable PCH (precompiled headers) in your build system
- Use ccache or similar (many other options that change regularly, Google for them)
- Be aware of unity builds
- Know the possibilities and limitations of extern template
- Use a build analysis tool to see where build time is spent

Use an IDE

This is the most surprising side effect of using a modern IDE that I have observed: IDE's do realtime analysis of the code. Realtime analysis means that you know as you are typing if the code is going to compile. Therefore, you spend less time waiting for builds.



Exercise: What are build times costing you?

Try to figure out how much build times are costing in developer time and see how much could be saved if build times were lessened.

Resources

- [A guide to unity builds¹](#)
- [Unity builds with Meson²](#)
- [Unity builds with CMake³](#)
- [PCH with Meson⁴](#)
- [PCH with CMake⁵](#)
- [ccache⁶](#)
- [CMake Compiler Launcher⁷](#)
- [Clang Build Analyzer⁸](#)
- [Getting started with C++ Build Insights⁹](#)
- [Introducing vcperf /timetrace for C++ build time analysis¹⁰](#)

¹<https://onqtam.com/programming/2018-07-07-unity-builds/>

²<https://mesonbuild.com/Unity-builds.html>

³https://cmake.org/cmake/help/latest/prop_tgt/UNITY_BUILD.html

⁴<https://mesonbuild.com/Precompiled-headers.html>

⁵https://cmake.org/cmake/help/latest/command/target_precompile_headers.html

⁶<https://ccache.dev/>

⁷https://cmake.org/cmake/help/latest/prop_tgt/LANG_COMPILER_LAUNCHER.html?highlight=ccache

⁸<https://github.com/aras-p/ClangBuildAnalyzer>

⁹<https://docs.microsoft.com/en-us/cpp/build-insights/get-started-with-cpp-build-insights?view=vs-2019>

¹⁰<https://devblogs.microsoft.com/cppblog/introducing-vcperf-timetrace-for-cpp-build-time-analysis/>

46. Use the Tools: Multiple Compilers

Support *at least* 2 compilers on your platform. Each compiler does different analyses and implements the standard in a slightly different way.

If you use Visual Studio, you should be able to switch between clang and cl.exe relatively easily. You can also use WSL and enable remote Linux Builds.

If you use Linux, you should be able to switch between GCC and Clang easily.



On MacOS, be sure the compiler you are using is what you think it is. The gcc command is likely a symlink to clang installed by Apple.

For installing newer or different compilers on your platform, the following is available:

Ubuntu / Debian

- GCC - [Toolchain PPA](#)¹
- Clang - [apt packages](#)²

Windows

- [GCC MinGW](#)³
- [Clang official downloads](#)⁴

¹<https://launchpad.net/~ubuntu-toolchain-r/+archive/ubuntu/ppa>

²<https://apt.llvm.org/>

³<http://mingw.org/>

⁴<https://releases.llvm.org/download.html>

MacOS

- Homebrew / MacPorts



Exercise: Add Another Compiler

Since you have already enabled continuous builds of your system, it's time to add another compiler.

A new version of the compiler you currently require is always a good idea. But if you only support GCC, consider adding Clang. Or if you only support Clang, add GCC. If you're on Windows, add MinGW or Clang in addition to MSVC.



Exercise: Add Another Operating System

Hopefully, at least some portion of your project can be ported to another operating system. The exercise of getting parts of the project compiling on another operating system and toolchain will teach you a lot about your code's nature.

Resources

- C++Now 2015: Jason Turner “Thinking Portable: How and Why to make your C++ Cross Platform”⁵

⁵<https://youtu.be/cb3WIL96N-o>

47. Fuzzing and Mutating

Your imagination limits the tests that you can create. Do you try to be malicious when calling your APIs? Do you intentionally pass malformed data to your inputs? Do you process inputs from unknown or untrusted sources?

Generating all possible inputs to all possible function calls in all possible combinations is impossible. Fortunately, tools exist to solve this problem.

Fuzzing

Fuzz testers generate strings of random data of various lengths. The test harness you write consumes these strings of data and processes them in some way that is appropriate for your application. The fuzz tester analyzes coverage data generated from your test's execution and uses that information to remove redundant tests and generate new novel and unique tests.

In theory, a fuzz test will eventually reach 100% code coverage of your tested code, if left to run long enough. Combined with AddressSanitizer, this makes a powerful tool for finding bugs in your code. [One interesting article from 2015¹](#) describes how the combination of a fuzz tester and AddressSanitizer could have found the security flaw “heartbleed” in OpenSSL in less than 6 hrs.



Fuzz testing primarily finds memory and security flaws.

Many different fuzzing tools exist. For the sake of this section, I am going to cover only LLVM's [libFuzzer²](#). All fuzz testers operate under the same premise.

¹<https://blog.hboeck.de/archives/868-How-Heartbleed-couldve-been-found.html>

²<https://www.llvm.org/docs/LibFuzzer.html>

You must provide some sort of entry point. The entry point generally takes the form of a function like:

libFuzzer entry point.

```
extern "C" int LLVMFuzzerTestOneInput(const uint8_t *Data,
                                     size_t Size);
```

The Data pointer is always valid, and the Size parameter is ≥ 0 .

If your library primarily parses input files (think libpng) then your job is quite easy:

libFuzzer data being used.

```
extern "C" int LLVMFuzzerTestOneInput(const uint8_t *Data,
                                     size_t Size)
{
    parseInput(Data, Size);
}
```

If your functions take data structures instead of input strings, your job is slightly more complicated but doable.

Advanced libFuzzer data usage.

```
template<typename Type>
std::pair<const uint8_t *, size_t, Type>
    createStruct(const uint8_t *Data, size_t Size)
{
    // we're only allowed to do this with trivial types
    static_assert(std::is_trivial_v<Type>);
    Type result{}; // default initialize
    const auto bytesToRead = std::min(sizeof(Type), Size);
    std::memcpy(&result, Data, bytesToRead);
    return {std::next(Data, bytesToRead), Size - bytesToRead, result};
}

extern "C" int LLVMFuzzerTestOneInput(const uint8_t *Data,
```

```
                                size_t Size)
{
    // This example is meant as inspiration, it has not been
    // tested in a real test
    auto [newDataPtr, remainingSize, Obj1]
        = createStruct<Type1>(Data,Size);
    auto [lastDataPtr, lastSize, Obj2]
        = createStruct<Type2>(newDataPtr, remainingSize);

    functionToTest(Obj1, Obj2);
}
```

The fuzzer will quickly learn that any new data input where `Size > sizeof(Type1) + sizeof(Type1)` does not generate new code paths and will focus on the appropriate amount of data.

Mutating

Mutation testing works by modifying conditionals and constants in the code being tested.

Pseudo code example.

```
bool greaterThanFive(const int value) {
    return value > 5; // comparison
}

void tests() {
    assert(greaterThanFive(6));
    assert(!greaterThanFive(4));
}
```

A mutation tester could modify the constant 5 or the > so the resulting code might become

Mutated code.

```
bool greaterThanFive(const int value) {  
    return value < 5; // mutated  
}
```

Any test that continues to pass is a “mutant that has survived” and may indicate either a flawed test or a bug in the code.



Exercise: Create a fuzz test harness.

Apply the examples demonstrated here to create fuzz testers for your code. What challenges do you hit?



Look at [FuzzedDataProvider.h](#)³ for more helper functions



Exercise: Investigate mutation testing.

The author of this book has no direct experience with mutation testing. Is it something you can use in your project? What interesting resources do you find?

Resources

- [C++Now 2018: Marshall Clow “Making Your Library More Reliable with Fuzzing”](#)⁴

³<https://github.com/llvm-mirror/compiler-rt/blob/master/include/fuzzer/FuzzedDataProvider.h>

⁴<https://youtu.be/LILJRHToyUk>

- C++ Weekly Ep 85: Fuzz Testing⁵
- CppCast: Alex Denisov “Mutation Testing With Mull”⁶
- NDC TechTown 2019: Seph De Busser “Testing The Tests: Mutation Testing for C++”⁷
- CppCon 2017: Kostya Serebryany “Fuzz or lose...”⁸
- CppCon 2020: Barnabás Bágyi “Fuzzing Class Interfaces for Generating and Running Tests with libFuzzer”⁹ - Inspirational talk about using fuzzing in novel ways. Video is not yet on YouTube, but look for it after this book is published.
- Autotest¹⁰ - Library associated with “Fuzzing Class Interfaces for Generating and Running Tests with libFuzzer” talk.

⁵<https://youtu.be/gO0KBoqkOoU>

⁶<https://cppcast.com/alex-denisov/>

⁷https://youtu.be/M-5_M8qZXaE

⁸<https://youtu.be/k-Cv8Q3zWNQ>

⁹<https://cppcon2020.sched.com/event/e7An/fuzzing-class-interfaces-for-generating-and-running-tests-with-libfuzzer?iframe=no>

¹⁰<https://gitlab.com/wilzegers/autotest/>

48. Continue Your C++ Education

You must continually learn if you want to become better at what you do, and many resources are available to you to continue your C++ education.

Know How To Ask Questions

Kate Gregory has published an [excellent article on how to ask questions](#)¹.

Some key points are:

- Don't use screenshots
- Use good variable names
- Add some tests
- Listen to what people are telling you

Conferences And Local User Groups

There is almost certainly one near you. It's a great way to network and learn new things. Check out the [ISO C++ Conferences Worldwide List](#)² and [Meeting C++'s User Groups List](#)³.

I am finishing this book during the global COVID-19 pandemic. So conferences and user groups are mostly on hold right now. But this presents an attractive new opportunity for many.

¹<http://www.gregcons.com/KateBlog/HowToAskForCCodingHelp.aspx>

²<https://isocpp.org/wiki/faq/conferences-worldwide>

³<https://meetingcpp.com/usergroups/>

Many of those conferences and user groups are now meeting online. It's now possible for us all to attend each other's user groups. The [North Denver Metro C++ Meetup](#)⁴, for example, regularly has one attendee from Thailand each month.



Note from the author: when interacting with the C++ community remember to treat others with dignity and respect; be patient; take time to understand the rules and norms of the particular community with which you are interacting.

C++ Weekly

This book references C++ Weekly throughout as a resource to go back to for more information and examples to share with your coworkers. At this moment, the show has been going for 235 weeks straight with many special editions, extras, and live streams.

cppreference.com

The website is fantastic, but you might not know that you can create an account and customize the content to the version of C++ you are using. Also, you can execute examples and [download an offline version](#)⁵!

Hire a Trainer to Come Onsite for Your Company

Team training gets your team thinking in a new direction, improves morale, and boosts employee retention. Since you made it this far, I'm going to offer you a coupon.

⁴<https://www.meetup.com/North-Denver-Metro-C-Meetup/>

⁵<https://en.cppreference.com/w/Cppreference>

If you mention this book, you'll get 10% off onsite training costs at your company from me. (travel costs not discounted). Hopefully, travel restrictions will not last much longer.

YouTube

- [Andreas Fertig's Channel](#)⁶
- [C++ Weekly \(Author's Channel\)](#)⁷
- [CopperSpice](#)⁸

⁶https://www.youtube.com/channel/UCxJflsPGHFS3_nRDv1u-Q8g

⁷<https://www.youtube.com/c/JasonTurner-lefticus>

⁸<https://www.youtube.com/copperspice>

49. Thank You

Sponsors

Thank you to all of my Book Supporter patrons who helped make this book possible!

Adam Albright
Adam P Shield
Alejandro Lucena
Alexander Roper
Andrei Sebastian Cîmpean
Anton Smyk
Arman Imani
Bill Baker
Björn Fahller
Brendan Nolan
Clint Rajaniemi
Corentin Gay
David C Black
Dennis Börm
Emyr Williams
Fedor Alekseev
Florian Sommer
Gwendolyn Hunt
Ivan Pakhomov
Jack Glass
Jaewon Jung
Jeff Bakst
Kacper Kołodziej
Lars Ove Larsen
Magnus Westin

Martin Hammerchmidt
Matt Godbolt
Matthew Guidry
Michael Pearce
Michael Pettit
Natalya Kochanova
Olafur Waage
Panos Gourgarris
Ralph Jeffrey Steinhagen
Reiner Eiteljoerge
Sebastian Raaphorst
Sergii Zaiets
Tim Butler
Tobias Dieterich
Tomasz Cwik
Yacob Cohen-Arazi

Reviewers

Craig Scott and Alexander Roper, thank you for extensive notes and feedback during prerelease.

50. Bonus: Understand The Lambda

A surprising complexity hides behind the simple lambda of C++. Initially added in C++11, it was initially constrained. With each version of C++, the lambda becomes more flexible and powerful.

Lambdas reverse some of the defaults from the rest of C++. Default `const` and automatically `constexpr` when possible; they give us some of what we wish the rest of the language could have.

Lambda grammar.

lambda-expression:

lambda-introducer lambda-declarator(opt) compound-statement

lambda-introducer < template-parameter-list > requires-clause(opt)
lambda-declarator(opt) compound-statement

lambda-introducer:

[lambda-capture(opt)]

lambda-declarator:

(parameter-declaration-clause) decl-specifier-seq(opt)
noexcept-specifier(opt) attribute-specifier-seq(opt)
trailing-return-type(opt) requires-clause(opt)

If you can [read standard-eze¹](http://eel.is/c++draft/expr.prim.lambda), you can dig into all of the features of C++20's lambdas yourself.

¹<http://eel.is/c++draft/expr.prim.lambda>

Allowed lambdas as of C++20.

```
// valid empty lambda, does nothing
[]{};
// optional to have parameter list
[](){};
// C++17 explicit constexpr and void return
[]() constexpr -> void {};
// immediately invoked lambda
auto i = [](){ return 42; }();
// Not allowed before C++17, because constexpr
constexpr auto j = []{ return 42; }();
// generic lambda, C++14
[](auto x){ return x + 42; };
// variadic lambda, C++14
[](auto ... x){ return std::vector<int>(x...); };
// capture by copy, C++11
[i](){ return i + 42; };
// generalized capture, C++14 (what's the type of i?)
[i = 42]{ return i + 42; };
// stateful lambda, C++11
[i]() mutable { return ++i; };
// explicit template, C++20
[]<typename T>(T x){ return x + 42; };

// C++14 generic lambda returning a C++20 lambda with variadic
// capture expression which returns a fold expression summation
// of the captured values.
[](auto ... val){ return [...val = val]{ return (val + ...); }; };
```

If you understand every aspect of C++'s lambdas and how the compiler implements them, you know everything important about C++.

This is why I put together my [C++ class on YouTube about lambdas](https://www.youtube.com/playlist?list=PLs3KjaCtOwSY_Awyliwm-fRjEOa-SRbs-)².

²https://www.youtube.com/playlist?list=PLs3KjaCtOwSY_Awyliwm-fRjEOa-SRbs-

In 2018 when compilers first started supporting C++20's new lambdas, I implemented this mostly standards-compliant version of `std::bind` using lambdas.

(Continued on next page.)

std::bind implemented with C++20 lambdas.

```
template <std::size_t Idx>
struct Placeholder {};
```

```
template <typename T>
struct Bound {
    constexpr decltype(auto) operator()(auto &&...param) const {
        return t(std::forward<decltype(param)>(param)...);
    }
}
```

```
    T t;
};
```

```
template <typename T>
Bound(T) -> Bound<T>;
```

```
template <std::size_t Idx, typename T>
constexpr decltype(auto) get_param(const Placeholder<Idx> &,
                                   T &&t) {
    return std::get<Idx>(t);
}
```

```
template <typename Param, typename T>
constexpr decltype(auto) get_param(Param &&param, T &&) {
    return std::forward<Param>(param);
}
```

```
template <typename Param, typename T>
constexpr decltype(auto) get_param(const Bound<Param> &b,
                                   T &&t) {
    return std::apply(b, std::forward<T>(t));
}
```

```
constexpr decltype(auto) bind(auto &&callable, auto &&...param) {
    return Bound{
```

```
[callable = std::forward<decltype(callable)>(callable),
... xs = std::forward<decltype(param)>(param)]
(auto &&...values) {
    auto passed_params =
        std::forward_as_tuple(
            std::forward<decltype(values)>(values)...);
    return std::invoke(callable,
                        get_param(xs, passed_params)...);
}
};
}
```

I haven't looked at this code in 2 years, but here is a Compiler Explorer link for you to play with.

<https://godbolt.org/z/hhde3P³>



Exercise: Understand the given example and critique it.

What should I have done differently with the above example? Can it be constrained with concepts? Does it need better names? What would you do differently?

³<https://godbolt.org/z/hhde3P>